

# Distributed Algorithms for Connectivity and MST in Large Graphs with Efficient Local Computation

Eric Ajieren, **Khalid Hourani**, William K. Moses Jr., Gopal  
Pandurangan

University of Houston

January 6, 2022

# Table of Contents

- 1 Introduction
- 2 Parallel Flooding
- 3 Borůvka-Style Algorithm
- 4 An Almost-Optimal Randomized Algorithm
- 5 Conclusion

# Table of Contents

- 1 Introduction
- 2 Parallel Flooding
- 3 Borůvka-Style Algorithm
- 4 An Almost-Optimal Randomized Algorithm
- 5 Conclusion

# The $k$ -Machine Model

- study distributed algorithms for large-scale graphs

# The $k$ -Machine Model

- study distributed algorithms for large-scale graphs
- focus on connectivity and MST

# The $k$ -Machine Model

- study distributed algorithms for large-scale graphs
- focus on connectivity and MST
- in  $k$ -machine model

# The $k$ -Machine Model

- study distributed algorithms for large-scale graphs
- focus on connectivity and MST
- in  $k$ -machine model
  - $k \geq 2$  machines jointly perform computations on input graph

# The $k$ -Machine Model

- study distributed algorithms for large-scale graphs
- focus on connectivity and MST
- in  $k$ -machine model
  - $k \geq 2$  machines jointly perform computations on input graph
  - input graph distributed uniformly at random to  $k$  machines



# The $k$ -Machine Model

- study distributed algorithms for large-scale graphs
- focus on connectivity and MST
- in  $k$ -machine model
  - $k \geq 2$  machines jointly perform computations on input graph
  - input graph distributed uniformly at random to  $k$  machines
    - $n$  nodes,  $m$  edges

# The $k$ -Machine Model

- study distributed algorithms for large-scale graphs
- focus on connectivity and MST
- in  $k$ -machine model
  - $k \geq 2$  machines jointly perform computations on input graph
  - input graph distributed uniformly at random to  $k$  machines
    - $n$  nodes,  $m$  edges
    - each **node** and its **adjacency list** is given to a random machine

# The $k$ -Machine Model

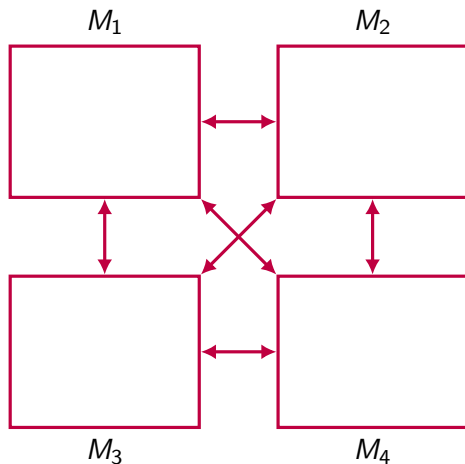
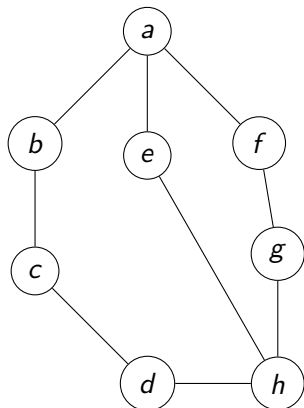
- study distributed algorithms for large-scale graphs
- focus on connectivity and MST
- in  $k$ -machine model
  - $k \geq 2$  machines jointly perform computations on input graph
  - input graph distributed uniformly at random to  $k$  machines
    - $n$  nodes,  $m$  edges
    - each **node** and its **adjacency list** is given to a random machine
    - commonly called **vertex centric** model

# The $k$ -Machine Model


- study distributed algorithms for large-scale graphs
- focus on connectivity and MST
- in  $k$ -machine model
  - $k \geq 2$  machines jointly perform computations on input graph
  - input graph distributed uniformly at random to  $k$  machines
    - $n$  nodes,  $m$  edges
    - each **node** and its **adjacency list** is given to a random machine
    - commonly called **vertex centric** model
  - assume  $n \gg k$  (e.g.  $k = \mathcal{O}(n^\epsilon)$ )

# $k$ -Machine Model

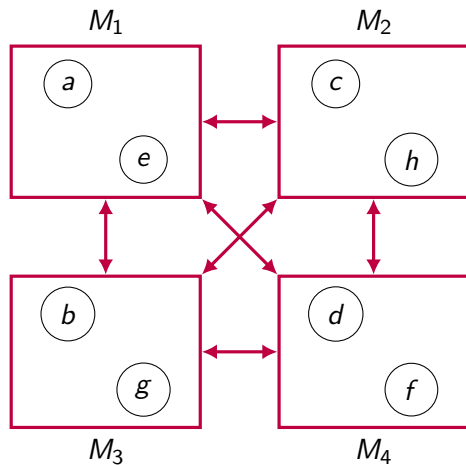
Input Graph



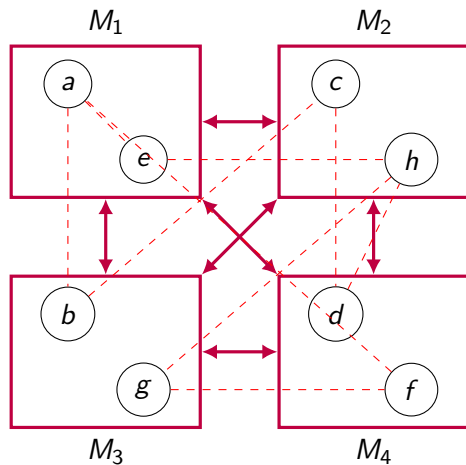
# $k$ -Machine Model



# $k$ -Machine Model



# $k$ -Machine Model





# $k$ -Machine Model

- typically, **only consider communication cost**

# $k$ -Machine Model

- typically, **only consider communication cost**
  - message complexity — number of communication rounds

# $k$ -Machine Model

- typically, **only consider communication cost**
  - message complexity — number of communication rounds
- traditionally done because network communication is much slower than local computation

# $k$ -Machine Model

- typically, **only consider communication cost**
  - message complexity — number of communication rounds
- traditionally done because network communication is much slower than local computation
  - less true as network speeds increase

# $k$ -Machine Model

- typically, **only consider communication cost**
  - message complexity — number of communication rounds
- traditionally done because network communication is much slower than local computation
  - less true as network speeds increase
  - for larger inputs, local computation becomes more significant

# $k$ -Machine Model

- typically, **only consider communication cost**
  - message complexity — number of communication rounds
- traditionally done because network communication is much slower than local computation
  - less true as network speeds increase
  - for larger inputs, local computation becomes more significant
- in practical implementations of  $k$ -machine model algorithms, often see suboptimal speedup

# $k$ -Machine Model

- typically, **only consider communication cost**
  - message complexity — number of communication rounds
- traditionally done because network communication is much slower than local computation
  - less true as network speeds increase
  - for larger inputs, local computation becomes more significant
- in practical implementations of  $k$ -machine model algorithms, often see suboptimal speedup
  - e.g. PageRank with message complexity  $\mathcal{O}(n/k)$  but whose wall-clock time is approximately  $n/k^{0.8}$

# Augmenting the $k$ -Machine Model with Local Computation

- we posit a new complexity measure



# Augmenting the $k$ -Machine Model with Local Computation

- we posit a new complexity measure
  - **local computation cost** —  $T_\ell$

# Augmenting the $k$ -Machine Model with Local Computation

- we posit a new complexity measure
  - **local computation cost** —  $T_\ell$
  - measures the worst-case local computation cost among  $k$  machines

# Augmenting the $k$ -Machine Model with Local Computation

- we posit a new complexity measure
  - **local computation cost** —  $T_\ell$
  - measures the worst-case local computation cost among  $k$  machines
- we refer to traditional communication complexity by  $T_c$

# Augmenting the $k$ -Machine Model with Local Computation

- we posit a new complexity measure
  - **local computation cost** —  $T_\ell$
  - measures the worst-case local computation cost among  $k$  machines
- we refer to traditional communication complexity by  $T_c$
- we analyze Connectivity and MST algorithms using this new complexity measure

# Summary of Results

- we analyze several MST/CC algorithms under the new complexity measure

# Summary of Results

- we analyze several MST/CC algorithms under the new complexity measure
- lower bounds for MST/CC

$$T_\ell = \Omega\left(\frac{m+n}{k} + \Delta + k\right)$$

$$T_c = \Omega\left(\frac{n}{k^2}\right)$$

# Summary of Results

- we analyze several MST/CC algorithms under the new complexity measure
- lower bounds for MST/CC

$$T_\ell = \Omega\left(\frac{m+n}{k} + \Delta + k\right)$$

$$T_c = \Omega\left(\frac{n}{k^2}\right)$$

Algorithm	Round complexity	Local runtime
Flooding	$\tilde{O}\left(\frac{n}{k} + D\right)$	$\tilde{O}\left(\frac{m}{k} + \Delta + k\right)$
• Filtering	$\tilde{O}\left(\frac{n}{k}\right)$	$\tilde{O}\left(\frac{m}{k} + n\right)$
Borůvka-Style	$\tilde{O}\left(\frac{n}{k}\right)$	$\tilde{O}\left(\frac{m+n}{k} + \Delta + k\right)$
Randomized CC	$\tilde{O}\left(\frac{n}{k^2}\right)$	$\tilde{O}\left(\frac{m+n}{k} + \Delta + k\right)$

# Some Useful Lemmas

- as a byproduct of our analysis, we have two useful lemmas for analysis of  $k$ -Machine Model algorithms



# Some Useful Lemmas

- as a byproduct of our analysis, we have two useful lemmas for analysis of  $k$ -Machine Model algorithms
  - the Node Distribution Lemma

# Some Useful Lemmas

- as a byproduct of our analysis, we have two useful lemmas for analysis of  $k$ -Machine Model algorithms
  - the Node Distribution Lemma
  - the Mapping Lemma

# Node Distribution Lemma

## Lemma (Node Distribution Lemma)

*Consider a graph  $G$  of nodes  $v_1, v_2, \dots, v_n$  with associated non-negative real-valued “weights”  $w(v)$  for each node  $v$ . Given a uniform, random distribution of the  $n$  nodes to  $k$  machines, as in the  $k$ -machine model, then, with probability at least  $1 - 1/n^a$  for any  $a > 0$ , the total weight of nodes at every machine is bounded above by*

$$\mathcal{O}(T_{\text{avg}} + \log n \cdot w_{\text{max}})$$

*where  $T_{\text{avg}} = \frac{1}{k} \sum_{i=1}^n w(v_i)$  and  $w_{\text{max}} = \max\{w(v_i)\}$ .*

# Mapping Lemma

## Lemma (Mapping Lemma)

*Let an  $n$ -node,  $m$ -edge graph  $G$  be partitioned among the  $k$  machines as  $N = \{p_1, \dots, p_k\}$ . Then with probability at least  $1 - 1/n^\alpha$ , where  $\alpha > 1$  is an arbitrary fixed constant, the following bounds hold:*

# Mapping Lemma

## Lemma (Mapping Lemma)

*Let an  $n$ -node,  $m$ -edge graph  $G$  be partitioned among the  $k$  machines as  $N = \{p_1, \dots, p_k\}$ . Then with probability at least  $1 - 1/n^\alpha$ , where  $\alpha > 1$  is an arbitrary fixed constant, the following bounds hold:*

- 1 *The number of vertices mapped to any machine is  $\mathcal{O}(n/k)$ .*

# Mapping Lemma

## Lemma (Mapping Lemma)

*Let an  $n$ -node,  $m$ -edge graph  $G$  be partitioned among the  $k$  machines as  $N = \{p_1, \dots, p_k\}$ . Then with probability at least  $1 - 1/n^\alpha$ , where  $\alpha > 1$  is an arbitrary fixed constant, the following bounds hold:*

- ① *The number of vertices mapped to any machine is  $\mathcal{O}(n/k)$ .*
- ② *The number of edges mapped to any machine is  $\mathcal{O}(m/k + \Delta \log n)$ .*

# Mapping Lemma

## Lemma (Mapping Lemma)

*Let an  $n$ -node,  $m$ -edge graph  $G$  be partitioned among the  $k$  machines as  $N = \{p_1, \dots, p_k\}$ . Then with probability at least  $1 - 1/n^\alpha$ , where  $\alpha > 1$  is an arbitrary fixed constant, the following bounds hold:*

- ① *The number of vertices mapped to any machine is  $\mathcal{O}(n/k)$ .*
- ② *The number of edges mapped to any machine is  $\mathcal{O}(m/k + \Delta \log n)$ .*
- ③ *The number of edges mapped to any link of the network is  $\mathcal{O}(m/k^2 + n/k)$ .*

# Table of Contents

- 1 Introduction
- 2 Parallel Flooding
- 3 Borůvka-Style Algorithm
- 4 An Almost-Optimal Randomized Algorithm
- 5 Conclusion



# Parallel Flooding

- the algorithm can be described from the perspective of a node

# Parallel Flooding

- the algorithm can be described from the perspective of a node
  - each node chooses an ID uniformly at random in  $[1, n^4]$

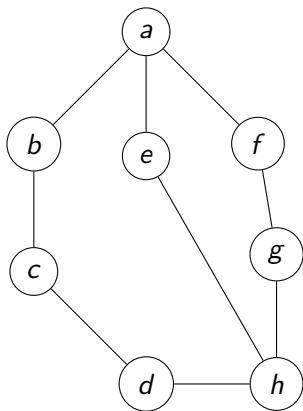
# Parallel Flooding

- the algorithm can be described from the perspective of a node
  - each node chooses an ID uniformly at random in  $[1, n^4]$
  - each node floods its ID

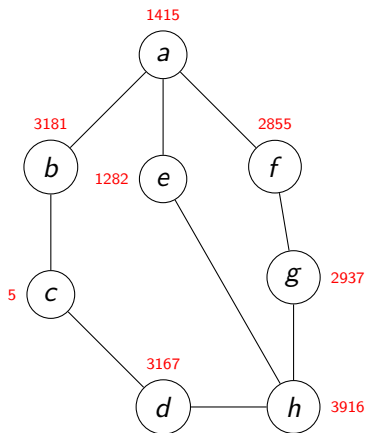
# Parallel Flooding

- the algorithm can be described from the perspective of a node
  - each node chooses an ID uniformly at random in  $[1, n^4]$
  - each node floods its ID
  - upon receiving a message (with an ID), if the new ID is greater than the current ID, the node updates its ID and floods


# Parallel Flooding



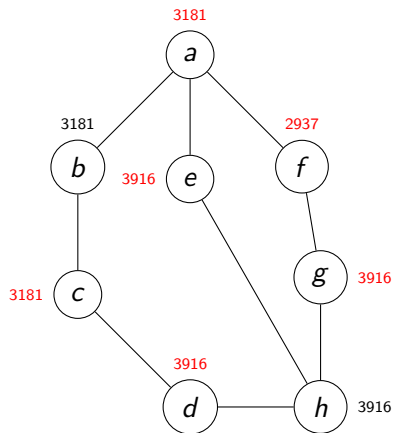
# Parallel Flooding



# Parallel Flooding

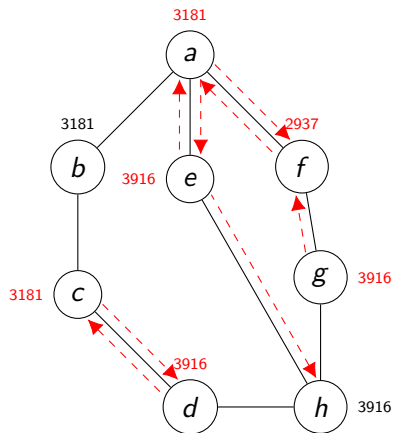


# Parallel Flooding

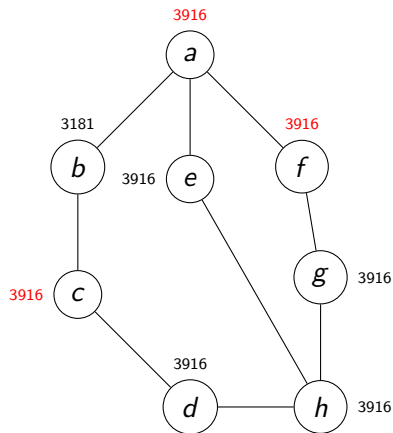




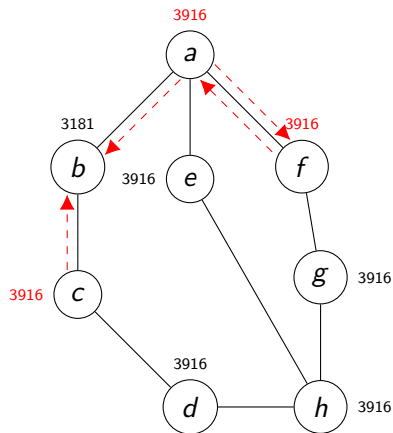
# Parallel Flooding



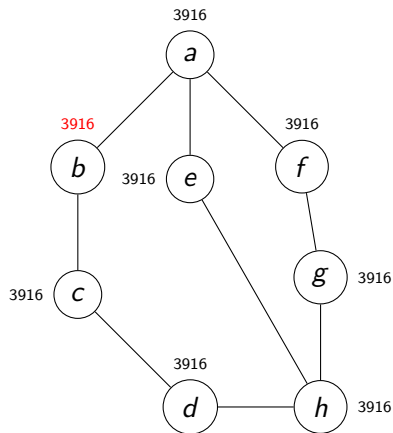
# Parallel Flooding



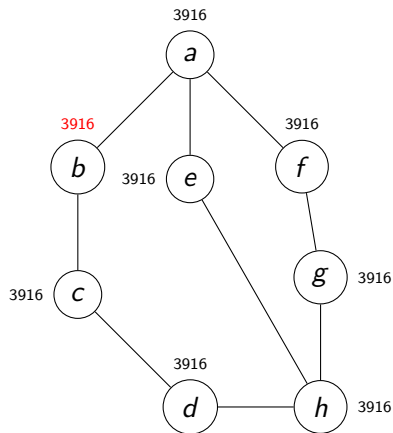
# Parallel Flooding



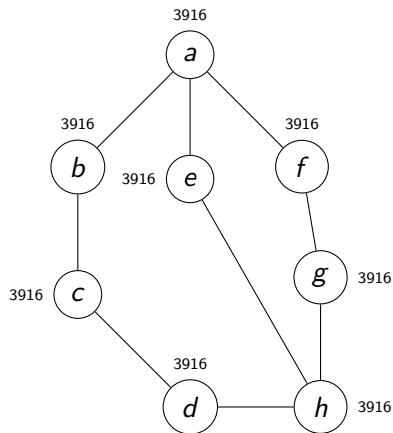
# Parallel Flooding



# Parallel Flooding



# Parallel Flooding



# Simulating in $k$ -Machine Model

- in order to simulate this in the  $k$ -Machine Model, each machine  $M$

# Simulating in $k$ -Machine Model

- in order to simulate this in the  $k$ -Machine Model, each machine  $M$ 
  - maintains list of nodes on the machine in descending-order of ID



# Simulating in $k$ -Machine Model

- in order to simulate this in the  $k$ -Machine Model, each machine  $M$ 
  - maintains list of nodes on the machine in descending-order of ID
  - iterates through list and simulates each node individually

# Simulating in $k$ -Machine Model

- in order to simulate this in the  $k$ -Machine Model, each machine  $M$ 
  - maintains list of nodes on the machine in descending-order of ID
  - iterates through list and simulates each node individually
    - when  $u$  located on  $M$  sends a message to  $v$  on  $M'$ ,  $M$  sends the appropriate message to  $M'$

# Simulating in $k$ -Machine Model

- in order to simulate this in the  $k$ -Machine Model, each machine  $M$ 
  - maintains list of nodes on the machine in descending-order of ID
  - iterates through list and simulates each node individually
    - when  $u$  located on  $M$  sends a message to  $v$  on  $M'$ ,  $M$  sends the appropriate message to  $M'$
  - after some  $\mathcal{O}(n \log n/k + D)$  rounds, aggregate all IDs located on machine  $M$  and send them to machine  $M_1$

# Simulating in $k$ -Machine Model

- in order to simulate this in the  $k$ -Machine Model, each machine  $M$ 
  - maintains list of nodes on the machine in descending-order of ID
  - iterates through list and simulates each node individually
    - when  $u$  located on  $M$  sends a message to  $v$  on  $M'$ ,  $M$  sends the appropriate message to  $M'$
  - after some  $\mathcal{O}(n \log n/k + D)$  rounds, aggregate all IDs located on machine  $M$  and send them to machine  $M_1$
- Finally, machine  $M_1$  counts the number of distinct IDs (which is the number of connected components) and broadcasts it to all machines

## Theorem

*With high probability, the above algorithm correctly counts the number of connected components with*

$$T_\ell = \mathcal{O}((m/k + \Delta \log n) \log n + k)$$

$$T_c = \mathcal{O}(n \log n/k + D)$$

## Theorem

*With high probability, the above algorithm correctly counts the number of connected components with*

$$T_\ell = \mathcal{O}((m/k + \Delta \log n) \log n + k)$$

$$T_c = \mathcal{O}(n \log n/k + D)$$

## Proof.



## Theorem

*With high probability, the above algorithm correctly counts the number of connected components with*

$$T_\ell = \mathcal{O}((m/k + \Delta \log n) \log n + k)$$

$$T_c = \mathcal{O}(n \log n / k + D)$$

## Proof.

- with Chernoff Bounds, can show that a node  $v$  will update max-ID  $\mathcal{O}(\log n)$  times with probability  $1 - 1/n^a$



## Theorem

*With high probability, the above algorithm correctly counts the number of connected components with*

$$T_\ell = \mathcal{O}((m/k + \Delta \log n) \log n + k)$$

$$T_c = \mathcal{O}(n \log n / k + D)$$

## Proof.

- with Chernoff Bounds, can show that a node  $v$  will update max-ID  $\mathcal{O}(\log n)$  times with probability  $1 - 1/n^a$
- $v$  will receive  $\mathcal{O}(\deg(v) \log n)$  higher IDs





## Theorem

*With high probability, the above algorithm correctly counts the number of connected components with*

$$T_\ell = \mathcal{O}((m/k + \Delta \log n) \log n + k)$$

$$T_c = \mathcal{O}(n \log n / k + D)$$

## Proof.

- with Chernoff Bounds, can show that a node  $v$  will update max-ID  $\mathcal{O}(\log n)$  times with probability  $1 - 1/n^a$
- $v$  will receive  $\mathcal{O}(\deg(v) \log n)$  higher IDs
- when max-ID is updated,  $v$  sends to  $\deg(v)$  nodes, totaling  $\mathcal{O}(\deg(v) \log n)$



## Theorem

*With high probability, the above algorithm correctly counts the number of connected components with*

$$T_\ell = \mathcal{O}((m/k + \Delta \log n) \log n + k)$$

$$T_c = \mathcal{O}(n \log n / k + D)$$

## Proof.

- with Chernoff Bounds, can show that a node  $v$  will update max-ID  $\mathcal{O}(\log n)$  times with probability  $1 - 1/n^a$
- $v$  will receive  $\mathcal{O}(\deg(v) \log n)$  higher IDs
- when max-ID is updated,  $v$  sends to  $\deg(v)$  nodes, totaling  $\mathcal{O}(\deg(v) \log n)$
- maximum runtime is therefore  $\mathcal{O}(\Delta \log n)$ .



Proof.



## Proof.

- apply node distribution lemma on the degree of nodes

$$\begin{aligned} T_\ell &= \mathcal{O}\left(\frac{1}{k}\left(\sum_{i=1}^n (d(v_i) \log n)\right) + \log n \cdot \Delta \log n\right) \\ &= \mathcal{O}\left(\left(\frac{m}{k} + \Delta \log n\right) \log n\right) \end{aligned}$$



## Proof.

- apply node distribution lemma on the degree of nodes

$$\begin{aligned} T_\ell &= \mathcal{O}\left(\frac{1}{k}\left(\sum_{i=1}^n (d(v_i) \log n)\right) + \log n \cdot \Delta \log n\right) \\ &= \mathcal{O}\left(\left(\frac{m}{k} + \Delta \log n\right) \log n\right) \end{aligned}$$

- in the worst case, a machine sends/receives messages from all other machines, totaling

$$\mathcal{O}\left(\left(\frac{m}{k} + \Delta \log n\right) \log n + k\right)$$



## Lemma

*The communication complexity of the algorithm is  $\mathcal{O}(n \log n/k + D)$  with high probability.*

## Lemma

*The communication complexity of the algorithm is  $\mathcal{O}(n \log n/k + D)$  with high probability.*

## Proof.



## Lemma

*The communication complexity of the algorithm is  $\mathcal{O}(n \log n/k + D)$  with high probability.*

## Proof.

- total number of rounds is  $\mathcal{O}(B \log n/(kW) + D)$  where  $W$  is the bandwidth,  $B$  is the number of broadcasts, and  $D$  is the diameter





## Lemma

*The communication complexity of the algorithm is  $\mathcal{O}(n \log n/k + D)$  with high probability.*

## Proof.

- total number of rounds is  $\mathcal{O}(B \log n/(kW) + D)$  where  $W$  is the bandwidth,  $B$  is the number of broadcasts, and  $D$  is the diameter
- taking  $W = \mathcal{O}(\log n)$



## Lemma

*The communication complexity of the algorithm is  $\mathcal{O}(n \log n/k + D)$  with high probability.*

## Proof.

- total number of rounds is  $\mathcal{O}(B \log n/(kW) + D)$  where  $W$  is the bandwidth,  $B$  is the number of broadcasts, and  $D$  is the diameter
- taking  $W = \mathcal{O}(\log n)$
- by mapping lemma,  $\mathcal{O}(\log n)$  broadcasts for a particular node with probability  $1 - 1/n^2$



## Lemma

*The communication complexity of the algorithm is  $\mathcal{O}(n \log n/k + D)$  with high probability.*

## Proof.

- total number of rounds is  $\mathcal{O}(B \log n/(kW) + D)$  where  $W$  is the bandwidth,  $B$  is the number of broadcasts, and  $D$  is the diameter
- taking  $W = \mathcal{O}(\log n)$
- by mapping lemma,  $\mathcal{O}(\log n)$  broadcasts for a particular node with probability  $1 - 1/n^2$
- by union bound,  $\mathcal{O}(\log n)$  broadcasts for all nodes with probability  $1 - 1/n$



## Lemma

*The communication complexity of the algorithm is  $\mathcal{O}(n \log n/k + D)$  with high probability.*

## Proof.

- total number of rounds is  $\mathcal{O}(B \log n/(kW) + D)$  where  $W$  is the bandwidth,  $B$  is the number of broadcasts, and  $D$  is the diameter
- taking  $W = \mathcal{O}(\log n)$
- by mapping lemma,  $\mathcal{O}(\log n)$  broadcasts for a particular node with probability  $1 - 1/n^2$
- by union bound,  $\mathcal{O}(\log n)$  broadcasts for all nodes with probability  $1 - 1/n$
- thus,  $\mathcal{O}(n \log n)$  broadcasts initiated with high probability



# Table of Contents

- 1 Introduction
- 2 Parallel Flooding
- 3 Borůvka-Style Algorithm**
- 4 An Almost-Optimal Randomized Algorithm
- 5 Conclusion

# Borůvka-Style Algorithm

- we present a deterministic algorithm similar to Borůvka's

# Borůvka-Style Algorithm

- we present a deterministic algorithm similar to Borůvka's
- a quick review of Borůvka's

# Borůvka-Style Algorithm

- we present a deterministic algorithm similar to Borůvka's
- a quick review of Borůvka's
  - each machine stores a disjoint-set (union-find)  $\mathcal{M}$



# Borůvka-Style Algorithm

- we present a deterministic algorithm similar to Borůvka's
- a quick review of Borůvka's
  - each machine stores a disjoint-set (union-find)  $\mathcal{M}$
  - algorithm works in “phases”

# Borůvka-Style Algorithm

- we present a deterministic algorithm similar to Borůvka's
- a quick review of Borůvka's
  - each machine stores a disjoint-set (union-find)  $\mathcal{M}$
  - algorithm works in “phases”
  - in each phase:

# Borůvka-Style Algorithm

- we present a deterministic algorithm similar to Borůvka's
- a quick review of Borůvka's
  - each machine stores a disjoint-set (union-find)  $\mathcal{M}$
  - algorithm works in “phases”
  - in each phase:
    - an MOE is found for each fragment

# Borůvka-Style Algorithm

- we present a deterministic algorithm similar to Borůvka's
- a quick review of Borůvka's
  - each machine stores a disjoint-set (union-find)  $\mathcal{M}$
  - algorithm works in “phases”
  - in each phase:
    - an MOE is found for each fragment
    - fragments are merged along MOEs

# Borůvka-Style Algorithm

- we present a deterministic algorithm similar to Borůvka's
- a quick review of Borůvka's
  - each machine stores a disjoint-set (union-find)  $\mathcal{M}$
  - algorithm works in “phases”
  - in each phase:
    - an MOE is found for each fragment
    - fragments are merged along MOEs
  - when exactly **one** fragment remains, it forms an MST

# The Algorithm

- rather than broadcasting MOEs, simulate unicast version (as in GHS)

# The Algorithm

- rather than broadcasting MOEs, simulate unicast version (as in GHS)
- additionally, we filter to reduce the number of edges in each machine to  $\mathcal{O}(n)$

# The Algorithm

- rather than broadcasting MOEs, simulate unicast version (as in GHS)
- additionally, we filter to reduce the number of edges in each machine to  $\mathcal{O}(n)$ 
  - machines create MSFs of their local subgraphs using Kruskal's



# The Algorithm

- rather than broadcasting MOEs, simulate unicast version (as in GHS)
- additionally, we filter to reduce the number of edges in each machine to  $\mathcal{O}(n)$ 
  - machines create MSFs of their local subgraphs using Kruskal's
  - “discard” edges that are not part of local MSF

# The Algorithm

- rather than broadcasting MOEs, simulate unicast version (as in GHS)
- additionally, we filter to reduce the number of edges in each machine to  $\mathcal{O}(n)$ 
  - machines create MSFs of their local subgraphs using Kruskal's
  - “discard” edges that are not part of local MSF
  - the remaining (at most  $k(n - 1)$  edges whp) are the only MST edges remaining

# The Algorithm

- rather than broadcasting MOEs, simulate unicast version (as in GHS)
- additionally, we filter to reduce the number of edges in each machine to  $\mathcal{O}(n)$ 
  - machines create MSFs of their local subgraphs using Kruskal's
  - “discard” edges that are not part of local MSF
  - the remaining (at most  $k(n - 1)$  edges whp) are the only MST edges remaining
- algorithm then continues on remaining subgraphs

# The Algorithm

- for a fixed machine  $M$  and fragment  $f$ ,  $M$  determines the **Local Minimum Outgoing Edge** (LOE) and broadcasts

# The Algorithm

- for a fixed machine  $M$  and fragment  $f$ ,  $M$  determines the **Local Minimum Outgoing Edge** (LOE) and broadcasts
- from the broadcast of the LOEs, each machine locally determines the global MOE for any fragment it contains

# The Algorithm

- for a fixed machine  $M$  and fragment  $f$ ,  $M$  determines the **Local Minimum Outgoing Edge** (LOE) and broadcasts
- from the broadcast of the LOEs, each machine locally determines the global MOE for any fragment it contains
- the machine then merges fragments

# Merging

- Note that we cannot merge all fragments at once, as that takes time proportional to the length of the fragment chain

# Merging

- Note that we cannot merge all fragments at once, as that takes time proportional to the length of the fragment chain
- use a technique similar to **controlled GHS**



# Merging

- Note that we cannot merge all fragments at once, as that takes time proportional to the length of the fragment chain
- use a technique similar to **controlled GHS**
  - create rooted tree  $F$  — each node is a fragment and there is an edge between nodes if they share an MOE

# Merging

- Note that we cannot merge all fragments at once, as that takes time proportional to the length of the fragment chain
- use a technique similar to **controlled GHS**
  - create rooted tree  $F$  — each node is a fragment and there is an edge between nodes if they share an MOE
  - construct maximal matching (using e.g. Cole-Vishkin)

# Merging

- Note that we cannot merge all fragments at once, as that takes time proportional to the length of the fragment chain
- use a technique similar to **controlled GHS**
  - create rooted tree  $F$  — each node is a fragment and there is an edge between nodes if they share an MOE
  - construct maximal matching (using e.g. Cole-Vishkin)
  - merge all matched edges and any edge where **exactly** one endpoint is matched

# $T_\ell$ and $T_c$

- $T_c = \tilde{O}(n/k)$

# $T_\ell$ and $T_c$

- $T_c = \tilde{O}(n/k)$
- $T_\ell = \tilde{O}(\frac{m+n}{k} + \Delta + k)$

# $T_\ell$ and $T_c$

- $T_c = \tilde{O}(n/k)$
- $T_\ell = \tilde{O}\left(\frac{m+n}{k} + \Delta + k\right)$ 
  - filtering requires  $\tilde{O}(m/k + \Delta)$

# $T_\ell$ and $T_c$

- $T_c = \tilde{O}(n/k)$
- $T_\ell = \tilde{O}\left(\frac{m+n}{k} + \Delta + k\right)$ 
  - filtering requires  $\tilde{O}(m/k + \Delta)$
  - matching requires  $\tilde{O}(n/k)$  rounds

# $T_\ell$ and $T_c$

- $T_c = \tilde{O}(n/k)$
- $T_\ell = \tilde{O}\left(\frac{m+n}{k} + \Delta + k\right)$ 
  - filtering requires  $\tilde{O}(m/k + \Delta)$
  - matching requires  $\tilde{O}(n/k)$  rounds
  - additional  $\tilde{O}(k)$  from a machine needing to process messages from  $k - 1$  other machines



# Table of Contents

- 1 Introduction
- 2 Parallel Flooding
- 3 Borůvka-Style Algorithm
- 4 An Almost-Optimal Randomized Algorithm**
- 5 Conclusion

# Sketch-Based Randomized Algorithm

- we analyze the 2018 SPAA algorithm of Pandurangan, Robinson, and Scquizzato

# Sketch-Based Randomized Algorithm

- we analyze the 2018 SPAA algorithm of Pandurangan, Robinson, and Scquizzato
- optimal up to polylog factors in terms of round complexity

# Sketch-Based Randomized Algorithm

- we analyze the 2018 SPAA algorithm of Pandurangan, Robinson, and Scquizzato
- optimal up to polylog factors in terms of round complexity
  - but  $\tilde{O}(n^2)$  local computation complexity

# Sketch-Based Randomized Algorithm

- we analyze the 2018 SPAA algorithm of Pandurangan, Robinson, and Scquizzato
- optimal up to polylog factors in terms of round complexity
  - but  $\tilde{O}(n^2)$  local computation complexity
- the algorithm is similar to Borůvka's algorithm, with the output being a **labeling** of the nodes, such that nodes in the same component have the same label

# Sketch-Based Randomized Algorithm

- we analyze the 2018 SPAA algorithm of Pandurangan, Robinson, and Scquizzato
- optimal up to polylog factors in terms of round complexity
  - but  $\tilde{O}(n^2)$  local computation complexity
- the algorithm is similar to Borůvka's algorithm, with the output being a **labeling** of the nodes, such that nodes in the same component have the same label
- relies on **linear graph sketches** to efficiently merge multiple components

# Sketches

- sketches are found by forming a vector of length  $\binom{n}{2}$  for each node  $v$  and an appropriate choice of an  $\mathcal{O}(\log^2 n) \times \binom{n}{2}$  matrix. This requires  $\tilde{\Omega}(n^2)$  local computation time.

# Sketches

- sketches are found by forming a vector of length  $\binom{n}{2}$  for each node  $v$  and an appropriate choice of an  $\mathcal{O}(\log^2 n) \times \binom{n}{2}$  matrix. This requires  $\tilde{\Omega}(n^2)$  local computation time.
- can be improved by using the 2015 sketches of King, Kutten, and Thorup



# Sketches

- sketches are found by forming a vector of length  $\binom{n}{2}$  for each node  $v$  and an appropriate choice of an  $\mathcal{O}(\log^2 n) \times \binom{n}{2}$  matrix. This requires  $\tilde{\Omega}(n^2)$  local computation time.
- can be improved by using the 2015 sketches of King, Kutten, and Thorup
- instead of storing vectors of length  $\binom{n}{2}$ , a leader machine broadcasts an **odd hash function**

# Sketches

- sketches are found by forming a vector of length  $\binom{n}{2}$  for each node  $v$  and an appropriate choice of an  $\mathcal{O}(\log^2 n) \times \binom{n}{2}$  matrix. This requires  $\tilde{\Omega}(n^2)$  local computation time.
- can be improved by using the 2015 sketches of King, Kutten, and Thorup
- instead of storing vectors of length  $\binom{n}{2}$ , a leader machine broadcasts an **odd hash function**
- each component uses this hash function to find an MOE

# Sketches

- sketches are found by forming a vector of length  $\binom{n}{2}$  for each node  $v$  and an appropriate choice of an  $\mathcal{O}(\log^2 n) \times \binom{n}{2}$  matrix. This requires  $\tilde{\Omega}(n^2)$  local computation time.
- can be improved by using the 2015 sketches of King, Kutten, and Thorup
- instead of storing vectors of length  $\binom{n}{2}$ , a leader machine broadcasts an **odd hash function**
- each component uses this hash function to find an MOE
- brings it to  $T_\ell = \tilde{O}((m+n)/k + \Delta + k)$

# Sketches

- sketches are found by forming a vector of length  $\binom{n}{2}$  for each node  $v$  and an appropriate choice of an  $\mathcal{O}(\log^2 n) \times \binom{n}{2}$  matrix. This requires  $\tilde{\Omega}(n^2)$  local computation time.
- can be improved by using the 2015 sketches of King, Kutten, and Thorup
- instead of storing vectors of length  $\binom{n}{2}$ , a leader machine broadcasts an **odd hash function**
- each component uses this hash function to find an MOE
- brings it to  $T_\ell = \tilde{O}((m+n)/k + \Delta + k)$ 
  - optimal, up to polylog factors!

# Table of Contents

- 1 Introduction
- 2 Parallel Flooding
- 3 Borůvka-Style Algorithm
- 4 An Almost-Optimal Randomized Algorithm
- 5 Conclusion

# Conclusion

- Augmented the  $k$ -machine model with a metric to measure local computation

# Conclusion

- Augmented the  $k$ -machine model with a metric to measure local computation
  - Many “optimal” algorithms are only optimal with respect to  $T_c$

# Conclusion

- Augmented the  $k$ -machine model with a metric to measure local computation
  - Many “optimal” algorithms are only optimal with respect to  $T_c$
- Analyzed well-known MIS/CC algorithms with respect to  $T_\ell$  and  $T_c$



# Conclusion

- Augmented the  $k$ -machine model with a metric to measure local computation
  - Many “optimal” algorithms are only optimal with respect to  $T_c$
- Analyzed well-known MIS/CC algorithms with respect to  $T_\ell$  and  $T_c$

# Conclusion

- Augmented the  $k$ -machine model with a metric to measure local computation
  - Many “optimal” algorithms are only optimal with respect to  $T_c$
- Analyzed well-known MIS/CC algorithms with respect to  $T_\ell$  and  $T_c$

Algorithm	Round complexity	Local runtime
Flooding	$\tilde{O}\left(\frac{n}{k} + D\right)$	$\tilde{O}\left(\frac{m}{k} + \Delta + k\right)$
Filtering	$\tilde{O}\left(\frac{n}{k}\right)$	$\tilde{O}\left(\frac{m}{k} + n\right)$
Borůvka-Style	$\tilde{O}\left(\frac{n}{k}\right)$	$\tilde{O}\left(\frac{m+n}{k} + \Delta + k\right)$
Randomized CC	$\tilde{O}\left(\frac{n}{k^2}\right)$	$\tilde{O}\left(\frac{m+n}{k} + \Delta + k\right)$

# Open Problems

- Analyze other graph problems in this model

# Open Problems

- Analyze other graph problems in this model
- Implement algorithms and compare wall-clock time with  $T_\ell$  and  $T_c$

# Open Problems

- Analyze other graph problems in this model
- Implement algorithms and compare wall-clock time with  $T_\ell$  and  $T_c$ 
  - e.g. are there constants  $\alpha$  and  $\beta$  (for a given set of hardware) such that wall-clock time is  $\alpha T_\ell + \beta T_c$ ?